

Maximum Likelihood Estimation (MLE) and Maximum a Posteriori (MAP)

Suppose you are given a dataset $\mathbb{D} = \{X, Y\}$, where

$$\begin{aligned} X &= [x^1 x^2 \dots x^N] \in \mathbb{R}^N \\ Y &= [y^1 y^2 \dots y^N] \in \mathbb{R}^N \end{aligned}$$

and $x^i \in \mathbb{R} \forall i = 1, \dots, N$, $y^i \in \mathbb{R}$, $\forall i = 1, \dots, N$. Moreover, suppose that:

$$y^i = \theta_1 + \theta_2 x^i + \dots + \theta_K (x^i)^{K-1} + e^i \quad \forall i = 1, \dots, N \quad (1)$$

Where $e^i \sim \mathcal{N}(0; \sigma^2)$ is random Gaussian Noise and $\theta = (\theta_1, \dots, \theta_n)^T$. It is known that the Maximum Likelihood Estimation (MLE) approach works by defining the conditional probability of y given x , $p_\theta(y|x)$ and then optimizes the parameters θ to maximize this probability distribution over \mathbb{D} . Moreover, it is also known that this approach can be made equivalent to the deterministic approach to solve such problems (the Least Square method) by taking the negative-log of $p_\theta(y|x)$. Indeed, by assuming that the noise e^i is Gaussian for each $i = 1, \dots, N$, we have

$$p_\theta(y|x) = \mathcal{N}(f_\theta(x^i), \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(f_\theta(x^i) - y)^2}$$

Thus

$$p_\theta(Y|X) = \prod_{i=1}^N p_\theta(y^i|x^i) \implies$$

$$\theta_{MLE} = \arg \max_{\theta \in \mathbb{R}^s} p_\theta(Y|X) = \arg \min_{\theta \in \mathbb{R}^s} -\log p_\theta(Y|X) = \arg \min_{\theta \in \mathbb{R}^s} -\log \prod_{i=1}^N p_\theta(y^i|x^i) = \arg \min_{\theta \in \mathbb{R}^s} \sum_{i=1}^N -\log p_\theta(y^i|x^i) =$$

$$\arg \min_{\theta \in \mathbb{R}^s} \sum_{i=1}^N \frac{1}{2\sigma^2} (f_\theta(x^i) - y^i)^2 = \arg \min_{\theta \in \mathbb{R}^s} \frac{1}{2\sigma^2} \|\mathbf{f}_\theta - \mathbf{y}\|_2^2$$

If $f_\theta(x) = \theta_1 + \theta_2 x^i + \dots + \theta_K (x^i)^{K-1} = \sum_{j=1}^K \phi_j(x^i) \theta_j$, with $\phi_j(x) = x^{j-1}$, we already shown that the problem above becomes

$$\theta_{MLE} = \arg \min_{\theta \in \mathbb{R}^s} \frac{1}{2\sigma^2} (\Psi(X)\theta - \mathbf{y})^2$$

where $\Psi(X)$ is the $N \times K$ matrix that contains the powers from 0 to $K-1$ of the input datapoints x^i as rows (the Vandermonde matrix associated with $(x^1, \dots, x^N)^T$), $\theta = (\theta_1, \dots, \theta_n)^T$ and $\mathbf{y} = (y^1, \dots, y^N)^T$.

Note that the above equation is equivalent to the function you optimized in the Exercise 3 of Lab3 with GD, with $A := \Phi(X)$, $x := \theta$ and $b := y$.

When it is unclear how to set the parameter K and it is impossible to use the error plot, it is required to use the Maximum A Posterior (MAP) approach.

To show how it works, suppose that we know that the parameters are normally distributed $\theta \sim \mathcal{N}(0; \sigma_\theta^2 I)$. Then we can use the Bayes Theorem to express the A Posteriori probability on y given x and θ as

$$p(\theta|X, Y) = \frac{p(Y|X, \theta)p(\theta)}{p(Y|X)}$$

The MAP solution searches for a set of parameters θ that maximizes $p(\theta|X, Y)$. Following the same reasoning as before,

$$\theta_{MAP} = \arg \max_{\theta} p(\theta|X, Y) = \arg \min_{\theta} \sum_{i=1}^N -\log p(\theta|x^i, y^i) = \arg \min_{\theta} \sum_{i=1}^N -\log p(y^i|x^i, \theta) - \log p(\theta)$$

Given the two optimization problem above, you are required to implement a program that compare the two solutions, that we will refer to as θ_{MLE} and θ_{MAP} .

1. Define a test problem in the following way:

- Let the user fix a positive integer $K > 0$, and define $\theta_{true} = (1, 1, \dots, 1)^T$ (you can also consider different θ_{true});
- Define an input dataset

$$X = [x^1 x^2 \dots x^N] \in \mathbb{R}^N$$

where the x^i are N uniformly distributed datapoints in the interval $[a, b]$, where $a < b$ are value that the user can select;

- Given a set of functions $\{\phi_1, \phi_2, \dots, \phi_K\}$, define the Generalized Vandermonde matrix $\Phi(X) \in \mathbb{R}^{N \times K}$, whose element in position i, j is $\phi_j(x^i)$. In particular, write a function defining the classical Vandermonde matrix where $\phi_j(x) = x^{j-1}$;
- Given a variance $\sigma^2 > 0$ defined by the user, compute

$$Y = \Phi(X)\theta_{true} + e$$

where $e \sim \mathcal{N}(0, \sigma^2 I)$ is Gaussian distributed noise with variance σ^2 . Note that the test problem defined in this way is very similar to what we did to define a test problem in the first Lab.

2. We now built a dataset $\mathbb{D} = \{(X, Y)\}$ such that $\theta_{true} = (1, 1, \dots, 1)^T \in \mathbb{R}^K$ is the **best** solution to the least squares problem

$$\Phi(X)\theta \approx Y$$

3. Pretend not to know the correct value of K . The first task is to try to guess it and use it to approximate the true solution θ_{true} by MLE and MAP. To do that:

- Write a function that takes as input the training data $\mathbb{D} = (X, Y)$ and K and returns the MLE solution (with Gaussian assumption) $\theta_{MLE} \in \mathbb{R}^K$ for that problem. Note that the loss function can be optimized by GD, SGD or Normal Equations.
- Write a function that takes as input a set of K -dimensional parameter vector θ and a test set $\mathcal{TE} = \{(X_{test}, Y_{test})\}$ and returns the average absolute error of the polynomial regressor $f_\theta(x)$ over X_{test} , computed as:

$$\frac{1}{N_{test}} \|f_\theta(X_{test}) - Y_{test}\|_2^2$$

- For different values of K , plot the training datapoints and the test datapoints with different colors, and visualize (as a continuous line) the learnt regression model $f_{\theta_{MLE}}(x)$. Comment the results.
- For increasing values of K , use the functions defined above to compute the training and test error, where the test set is generated by sampling N_{test} new points on the same interval $[a, b]$ of the training set and generating the corresponding Y_{test} with the same procedure of the training set. Plot the two errors with respect to K . Comment the results.
- Write a function that takes as input the training data $\mathbb{D} = (X, Y)$, K and $\lambda > 0$ and returns the MAP solution (with Gaussian assumption) $\theta_{MAP} \in \mathbb{R}^K$ for that problem. Note that the loss function can be optimized by GD, SGD or Normal Equations.
- For K lower, equal and greater than the correct degree of the test polynomial, plot the training datapoints and the test datapoints with different colors, and visualize (as a continuous line) the learnt regression model $f_{\theta_{MAP}}(x)$ with different values of λ . Comment the results.
- For K being way greater than the correct degree of the polynomial, compute the MLE and MAP solution. Compare the test error of the two, for different values of λ (in the case of MAP).
- For K greater than the true degree of the polynomial, define $Err(\theta) = \frac{\|\theta - \theta_{true}\|_2}{\|\theta_{true}\|_2}$ where θ_{true} has been padded with zeros to match the shape of θ . Compute $Err(\theta_{MLE})$ and $Err(\theta_{MAP})$ for increasing values of K and different values of λ .
- Compare the results obtained by increasing the number N of datapoints.
- Compare the results obtained by the three algorithms GD, SGD and Normal Equations.

4. Repeat the first point of the homework by setting

$$y^i \sim Poi(y|\theta_{true,1} + \theta_{true,2}x^i + \dots + \theta_{true,K}(x^i)^{K-1}) \quad \forall i = 1, \dots, N \quad (2)$$

Where $Poi(y|\lambda)$ is the Poisson distribution with parameter λ (you can find the Python function to do that), such that

$$Poi(y|\lambda) = \frac{\lambda^y e^{-\lambda}}{y!}$$

And consequently,

$$p_{\theta}(y|x) = Poi(y|\theta_1 + \theta_2 x^i + \dots + \theta_K (x^i)^{K-1})$$

5. Compute the MLE and MAP losses for this choice of $p_{\theta}(y|x)$ and repeat the experiments in the previous points.